

Last course: Intensity Level Resolution



Low Detail



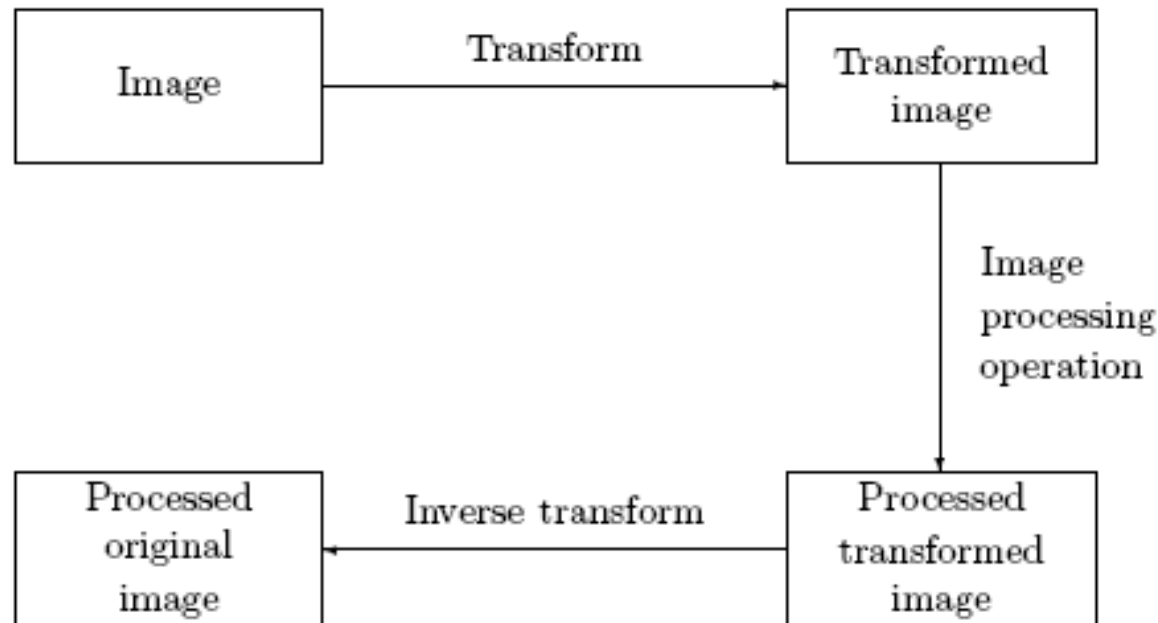
Medium Detail



High Detail

Last course POINT PROCESSING

Transform: represents the pixel values in some other, but equivalent form




Schema for transform processing

Neighbourhood processing. To change the grey level of a given pixel we need only know the value of the grey levels in a small neighbourhood of pixels around the given pixel.

Point operations. A pixel's grey value is changed without any knowledge of its surrounds.

Last Course Basic Transformations

FIGURE 3.3 Some basic gray-level transformation functions used for image enhancement.



Negative:

$$s = L - 1 - r$$

Log:

Inverse Log:

$$s = e^{cr} - 1$$

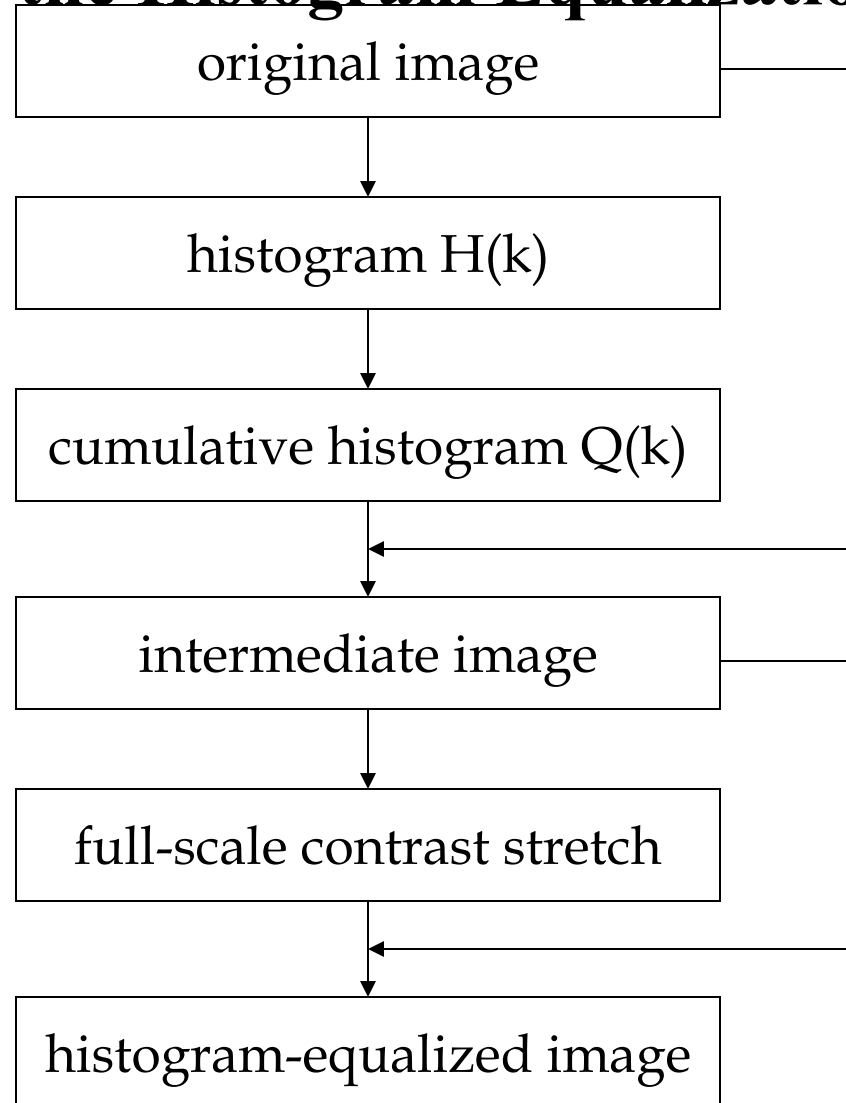
Power-law:

$$s = cr^\gamma$$

.....

Last course

Summary of the Histogram Equalization Algorithm





SPATIAL FILTERING

- Understand the mechanics of spatial filtering, and how spatial filters are formed
- Understand the principles of spatial convolution and correlation.
- Be familiar with the principal types of spatial filters, and how they are applied.
- Be aware of the relationships between spatial filters, and the fundamental role of lowpass filters.
- Understand how to use combinations of enhancement methods in cases where a single approach is insufficient.



$$g(x,y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s,t)f(x+s,y+t)$$



Linear image processing

- Image processing system $S(\cdot)$ is linear, iff superposition principle holds:

$$S(\alpha \cdot f[x,y] + \beta \cdot g[x,y]) = \alpha \cdot S(f[x,y]) + \beta \cdot S(g[x,y]) \quad \text{for all } \alpha, \beta \in \mathbb{R}$$

- Any linear image processing system can be written as


$$\vec{g} = H\vec{f}$$

Note: matrix H need not be square.

by sorting pixels into a column vector

$$\vec{f} = \begin{pmatrix} f[0,0] & f[1,0] & \cdots & f[N-1,0] & f[0,1] & \cdots & f[N-1,1] & \cdots & f[0,L-1] & \cdots & f[N-1,L-1] \end{pmatrix}^T$$

Convolution (1D)



Convolution (1D)


1	2	1
---	---	---

1	1	2	2	1	1	2	2	1	1
---	---	---	---	---	---	---	---	---	---




	$\frac{5}{4}$								
--	---------------	--	--	--	--	--	--	--	--


Convolution (1D)




Convolution (1D)




Convolution (1D)




Convolution (1D)



Convolution (1D)



Convolution (1D)




Convolution (1D)

This process is called
Convolution!!

1	2	1
---	---	---

1	1	2	2	1	1	2	2	1	1
---	---	---	---	---	---	---	---	---	---



	$\frac{5}{4}$	$\frac{7}{4}$	$\frac{7}{4}$	$\frac{5}{4}$	$\frac{5}{4}$	$\frac{7}{4}$	$\frac{7}{4}$	$\frac{5}{4}$	
--	---------------	---------------	---------------	---------------	---------------	---------------	---------------	---------------	--

Convolution/correlation on images

- The filter is now 2D
- Kernel (mask), kernel coefficients
- Size: **3x3**, 5x5, 7x7,

Normalisation

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

Input

1	2	0	1	3	
2	1	4	2	2	
1	0	1	0	1	
1	2	1	0	2	
2	5	3	1	2	

Output

	$\frac{12}{9}$				

Convolution/correlation on images

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

Input

Output

1	2	0	1	3	
2	1	4	2	2	
1	0	1	0	1	
1	2	1	0	2	
2	5	3	1	2	

	$\frac{12}{9}$	$\frac{11}{9}$			

Convolution/correlation on images

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

Input

1	2	0	1	3	
2	1	4	2	2	
1	0	1	0	1	
1	2	1	0	2	
2	5	3	1	2	

Output

	$\frac{12}{9}$	$\frac{11}{9}$	$\frac{14}{9}$		

Convolution/correlation on images

Input

 $\frac{1}{9}$

1	1	1
1	1	1
1	1	1

Output

1	2	0	1	3	
2	1	4	2	2	
1	0	1	0	1	
1	2	1	0	2	
2	5	3	1	2	

	$\frac{12}{9}$	$\frac{11}{9}$	$\frac{14}{9}$		
	$\frac{13}{9}$	$\frac{11}{9}$	$\frac{13}{9}$		
	$\frac{16}{9}$	$\frac{12}{9}$	$\frac{11}{9}$		

Math. of 2D Convolution/Correlation

Convolution

$$g(x, y) = h * f(x, y) = \sum_{j=-n}^n \sum_{i=-m}^m h(i, j) f(x-i, y-j)$$

Correlation


$$g(x, y) = h \circ f(x, y) = \sum_{j=-n}^n \sum_{i=-m}^m h(i, j) f(x+i, y+j)$$

Note: When the filter is symmetric: correlation = convolution!

1	1	1
1	1	1
1	1	1

2	3	2
-1	0	-1
2	3	2

Spatial Filtering



104	100	108
99	106	98
95	90	85

Original Image
Pixels

$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$

Filter

$$\begin{aligned}
 e &= \frac{1}{9} * 106 + \\
 &\quad \frac{1}{9} * 104 + \frac{1}{9} * 100 + \frac{1}{9} * 108 + \\
 &\quad \frac{1}{9} * 99 + \frac{1}{9} * 98 + \\
 &\quad \frac{1}{9} * 95 + \frac{1}{9} * 90 + \frac{1}{9} * 85 \\
 &= 98.3333
 \end{aligned}$$

The above is repeated for every pixel in the original image.

$$R = w_1 z_1 + w_2 z_2 + \dots + w_{mn} z_{mn}$$

$$= \sum_{k=1}^{mn} w_k z_k$$

$$= \mathbf{w}^T \mathbf{z}$$

Problems at the borders

- Why is the output image smaller than the input?
We are lacking information
- The bigger the kernel the bigger the problem
- Does it matter? Yes, if we are going to combine the images afterwards

Input


1	2	0	1	3	
2	1	4	2	2	
1	0	1	0	1	
1	2	1	0	2	
2	5	3	1	2	

Output

	$\frac{12}{9}$	$\frac{11}{9}$	$\frac{14}{9}$		
	$\frac{13}{9}$	$\frac{11}{9}$	$\frac{13}{9}$		
	$\frac{16}{9}$	$\frac{12}{9}$	$\frac{11}{9}$		
					23


Problems at the borders

What value should these
outside pixels have?





Strange Things Happen At The Edges!

At the edges of an image we are missing pixels to form a neighbourhood




Problems at the borders

- Solutions:
 - Complex and perhaps wrong



Smoothing Spatial Filters



a b

FIGURE 3.32 Two 3×3 smoothing (averaging) filter masks. The constant multiplier in front of each mask is equal to 1 divided by the sum of the values of its coefficients, as is required to compute an average.

Image Smoothing Example

- The image at the top left is an original image of size 500*500 pixels
- The subsequent images show the image after filtering with an averaging filter of increasing sizes
 - 3, 5, 9, 15 and 35
- Notice how detail begins to disappear




Image Smoothing Example




Image Smoothing Example



Image Smoothing Example



Image Smoothing Example




Image Smoothing Example






Image Smoothing Example




Another Smoothing Example

- By smoothing the original image we get rid of lots of the finer detail which leaves only the gross features for thresholding



Original Image



Smoothed Image



Thresholded Image

Convolution examples



Original
Bike



Bike blurred by convolution
Impulse response „box filter“

$$\frac{1}{25} \begin{pmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & [1] & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{pmatrix}$$

Convolution examples



Original
Bike



Bike blurred horizontally
Filter impulse response

$$\frac{1}{5} \begin{pmatrix} 1 & 1 & [1] & 1 & 1 \end{pmatrix}$$

Convolution examples




Original
Bike



Bike blurred vertically
Filter impulse response

$$\frac{1}{5} \begin{pmatrix} 1 \\ 1 \\ [1] \\ 1 \\ 1 \end{pmatrix}$$

Practice with linear filters




Original

0	0	0
0	1	0
0	0	0


?

Practice with linear filters




Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

Practice with linear filters




Original

0	0	0
0	0	1
0	0	0


?

Practice with linear filters




Original

0	0	0
0	0	1
0	0	0



Shifted left
By 1 pixel

Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

−


$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

?

(Note that filter sums to 1)

Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

−

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1




Sharpening filter

- Accentuates differences with local average

Order-Statistic (Nonlinear) Filters

			5	8	10	15	20	22	25	32	100
10	22	20									
5	8	15			10	22	20				
32	25	100			5	20	15				
					32	25	100				



a b c

FIGURE 3.35 (a) X-ray image of circuit board corrupted by salt-and-pepper noise. (b) Noise reduction with a 3×3 averaging mask. (c) Noise reduction with a 3×3 median filter. (Original image courtesy of Mr. Joseph E. Pascente, Lixi, Inc.)

Sharpening Spatial Filters

Previously we have looked at smoothing filters which remove fine detail

Sharpening spatial filters seek to highlight fine detail


- Remove blurring from images
- Highlight edges

Sharpening filters are based on *spatial differentiation*



Spatial Differentiation

Differentiation measures the *rate of change* of a function

Let's consider a simple 1 dimensional example



Spatial Differentiation



1st Derivative

The formula for the 1st derivative of a function is as follows:

$$\frac{\partial f}{\partial x} = f(x+1) - f(x)$$

It's just the difference between subsequent values and measures the rate of change of the function

2nd Derivative

The formula for the 2nd derivative of a function is as follows:

$$\frac{\partial^2 f}{\partial^2 x} = f(x+1) + f(x-1) - 2f(x)$$

Simply takes into account the values both before and after the current value

Using Second Derivatives For Image Enhancement

The 2nd derivative is more useful for image enhancement than the 1st derivative

- Stronger response to fine detail
- Simpler implementation
- We will come back to the 1st order derivative later on

The first sharpening filter we will look at is the *Laplacian*

- Isotropic
- One of the simplest sharpening filters
- We will look at a digital implementation

The Laplacian

The Laplacian is defined as follows:

$$\nabla^2 f = \frac{\partial^2 f}{\partial^2 x} + \frac{\partial^2 f}{\partial^2 y}$$

where the partial 1st order derivative in the x direction is defined as follows:

$$\frac{\partial^2 f}{\partial^2 x} = f(x+1, y) + f(x-1, y) - 2f(x, y)$$

and in the y direction as follows:

$$\frac{\partial^2 f}{\partial^2 y} = f(x, y+1) + f(x, y-1) - 2f(x, y)$$

The Laplacian (cont...)

So, the Laplacian can be given as follows:

$$\nabla^2 f = [f(x+1, y) + f(x-1, y) + f(x, y+1) + f(x, y-1)] - 4f(x, y)$$

We can easily build a filter based on this


0	1	0
1	-4	1
0	1	0

The Laplacian (cont...)


Applying the Laplacian to an image we get a new image that highlights edges and other discontinuities



Original
Image



Laplacian
Filtered Image



Laplacian
Filtered Image
Scaled for Display

But That Is Not Very Enhanced!

The result of a Laplacian filtering is not an enhanced image

We have to do more work in order to get our final image

Subtract the Laplacian result from the original image to generate our final sharpened enhanced image

$$g(x, y) = f(x, y) - \nabla^2 f$$




Laplacian
Filtered Image
Scaled for Display

Laplacian Image Enhancement




Original
Image

-



Laplacian
Filtered Image

=



Sharpened
Image

In the final sharpened image edges and fine detail are much more obvious

Laplacian Image Enhancement




Simplified Image Enhancement

The entire enhancement can be combined into a single filtering operation



$$\begin{aligned} g(x, y) &= f(x, y) - \nabla^2 f \\ &= f(x, y) - [f(x+1, y) + f(x-1, y) \\ &\quad + f(x, y+1) + f(x, y-1) \\ &\quad - 4f(x, y)] \\ &= 5f(x, y) - f(x+1, y) - f(x-1, y) \\ &\quad - f(x, y+1) - f(x, y-1) \end{aligned}$$

Simplified Image Enhancement (cont...)

This gives us a new filter which does the whole job for us in one step



Simplified Image Enhancement (cont...)



Variants On The Simple Laplacian

$$g(x, y) = f(x, y) + c[\nabla^2 f(x, y)]$$




There are lots of slightly different versions of the Laplacian that can be used:

0	1	0
1	-4	1
0	1	0

Simple
Laplacian

1	1	1
1	-8	1
1	1	1

Variant of
Laplacian



Unsharp Masking and Highboost Filtering

$$g_{\text{mask}}(x, y) = f(x, y) - \bar{f}(x, y)$$

$$g(x, y) = f(x, y) + k g_{\text{mask}}(x, y)$$



Original image.



Result of blurring with a Gaussian filter.



Unsharp mask.



Result of using unsharp masking.



Result of using highboost filtering.

Sobel Operators




Based on the previous equations we can derive the *Sobel Operators*

-1	-2	-1
0	0	0
1	2	1

-1	0	1
-2	0	2
-1	0	1

To filter an image it is filtered using both operators the results of which are added together

Sobel Example



An image of a contact lens which is enhanced in order to make defects (at four and five o'clock in the image) more obvious

Sobel filters are typically used for edge detection

1st & 2nd Derivatives

Comparing the 1st and 2nd derivatives we can conclude the following:


- 1st order derivatives generally produce thicker edges
- 2nd order derivatives have a stronger response to fine detail e.g. thin lines
- 1st order derivatives have stronger response to grey level step
- 2nd order derivatives produce a double response at step changes in grey level

Combining Spatial Enhancement Methods


Successful image enhancement is typically not achieved using a single operation

Rather we combine a range of techniques in order to achieve a final result

This example will focus on enhancing the bone scan to the right



Combining Spatial Enhancement Methods (cont...)




(a)

Laplacian filter of
bone scan (a)




(b)

Sharpened version of
bone scan achieved
by subtracting (a)
and (b)



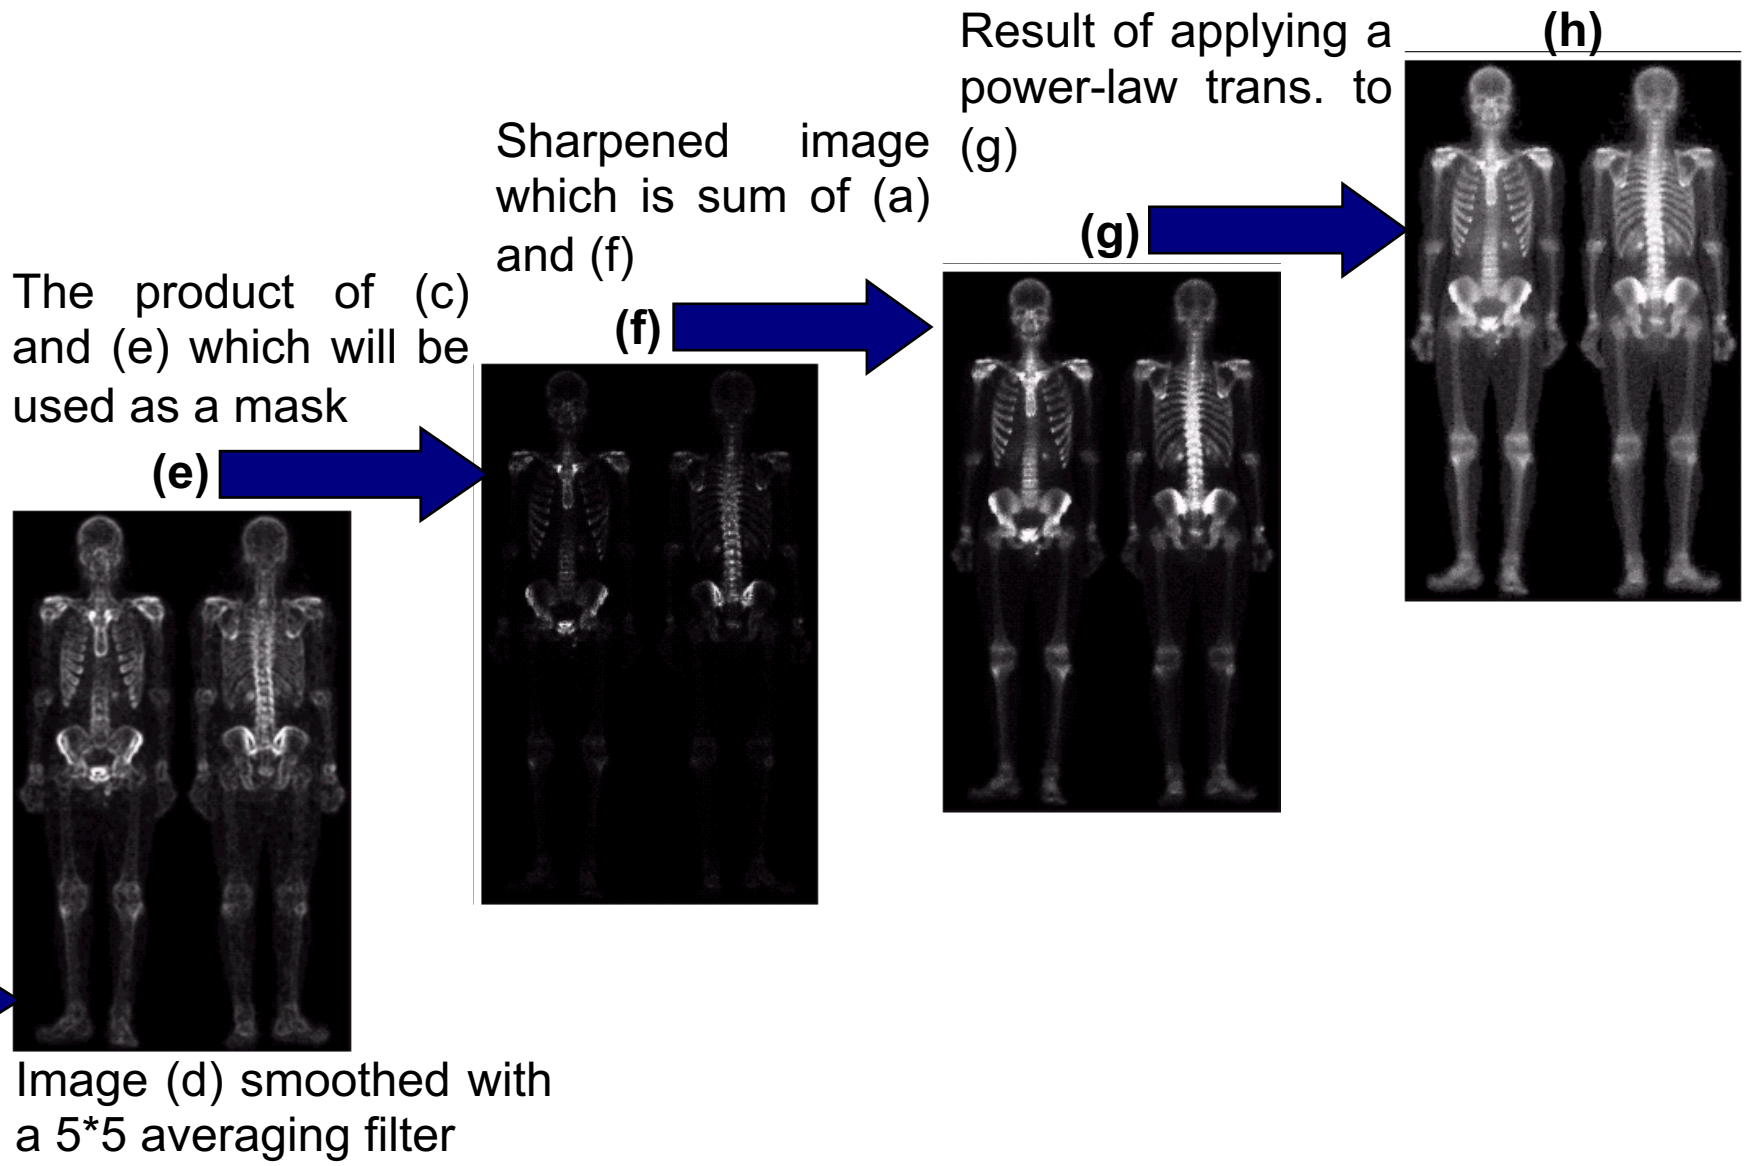
(c)

Sobel filter of bone
scan (a)



(d)

Combining Spatial Enhancement Methods (cont...)




Combining Spatial Enhancement Methods (cont...)


Compare the original and final images




Next Course: Frequency Spectra



=



+



=

